WIENER AI SOFTWARE STAMMTISCH

by Telos *Circle*

WAISS



#1 10.11.2026 · 18:00 CAFE RITTER

Telos *Circle*

Vienna-based nonprofit fostering interdisciplinary innovation across science, technology and entrepreneurship. Home of the WAISS.

HACKATHONS

OpenClaw HACK_001

HACK_002



Peter Wimberger

@TOPSREK

PROJECTS

Intercortex, Machine Literacy

LINKEDIN

/in/pwimberger

AGENDA

Tonight

01	AI recap of the month	20 min
02	Show & Tell · short talks	[xx min]
03	Open round & questions	open end



SLIDES & WEBSITE

waiss.teloscircle.com

Tonight’s slides, past and upcoming editions.

Want to present something next time?

Open a pull request on github.com/topsrek/waiss

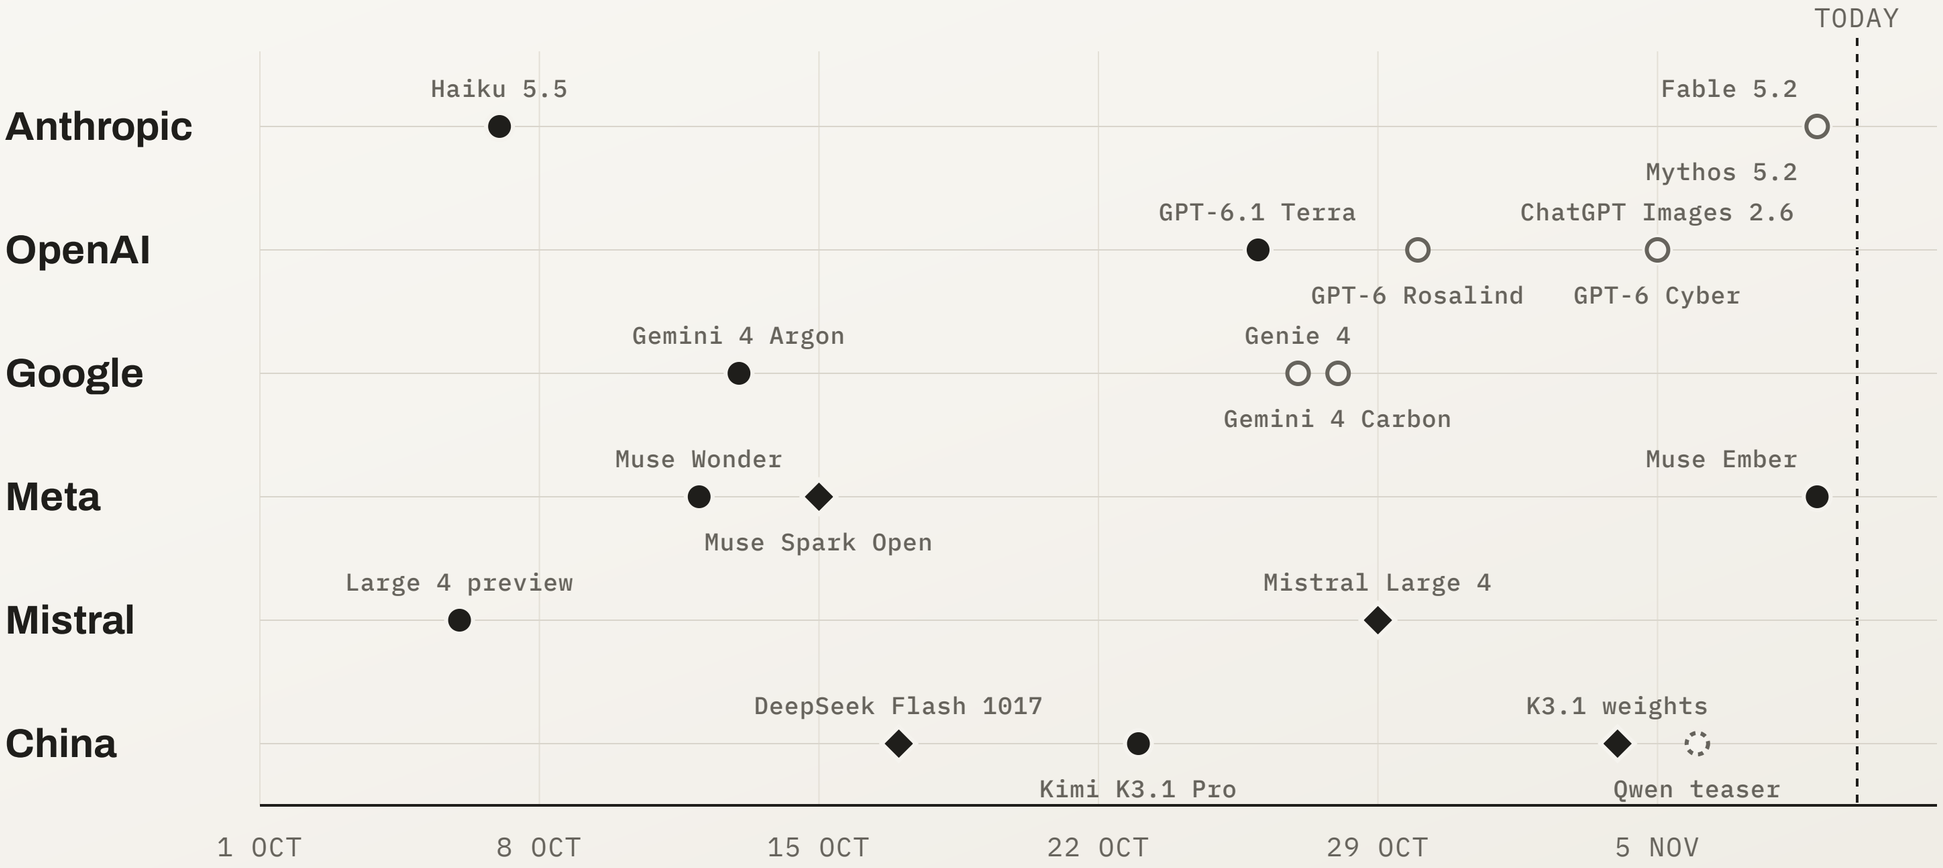
or email events@teloscircle.com

HTML PREFERRED · PPT AND PDF WORK TOO



RELEASES · 1 OCT – 10 NOV 2026

Eight labs, six weeks



INK DOT = AVAILABLE TO EVERYONE · DIAMOND = OPEN WEIGHTS · HOLLOW = SELECTED CUSTOMERS ONLY (CYBER AND LIFE-SCIENCE PROGRAMS, TRUSTED TESTERS) · DASHED = ANNOUNCED ONLY. ARGON: DATE OF GENERAL AVAILABILITY.



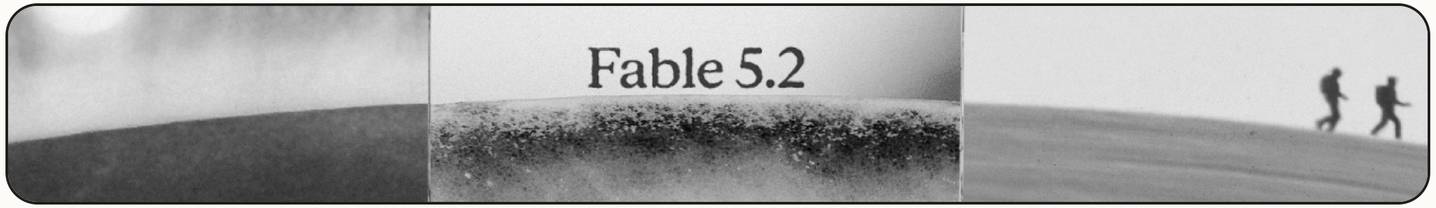
by Telos *Circle*

RELEASES · ANTHROPIC

Fable 5.2

Claude
@claudeai

Introducing Claude Fable 5.2 and Claude Mythos 5.2, designed to tackle the most daunting technical challenges.



19:57 · 9 Nov 2026



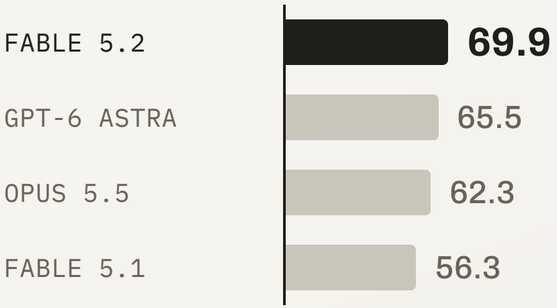
RELEASES · ANTHROPIC

Fable 5.2

Ahead almost everywhere, 20% cheaper than 5.1

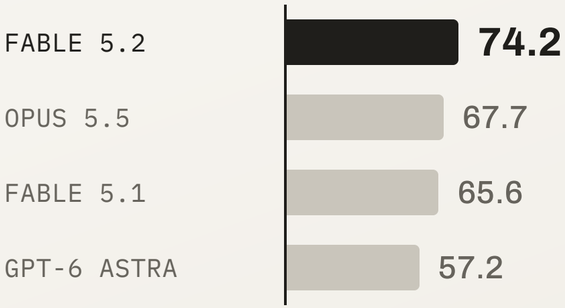
FrontierSWE v2

Agentic coding



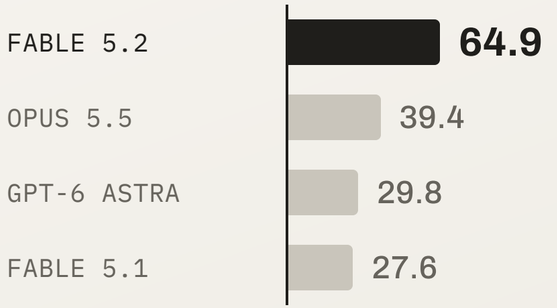
Humanity's Last Exam

Expert questions, with tools



RoboDojo Sim

Robot control in simulation



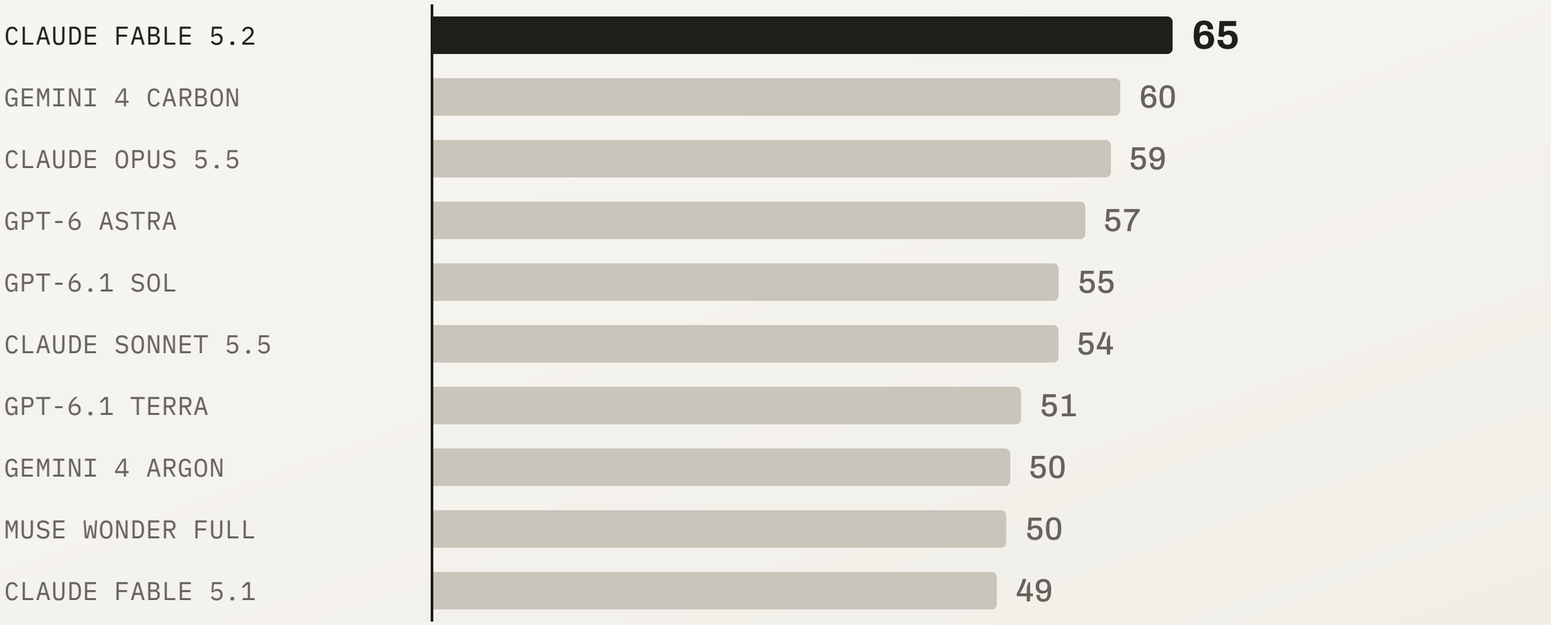


SCORE IN %, VENDOR-REPORTED, HIGHEST EFFORT PER MODEL. GEMINI 4 CARBON LEADS AUTOMATIONBENCH (53.9 VS 51.4). SOURCE: ANTHROPIC.COM/CLAUDE-FABLE-AND-MYTHOS-5-2

BENCHMARKS · ARTIFICIAL ANALYSIS

Fable 5.2

New #1 on the Intelligence Index



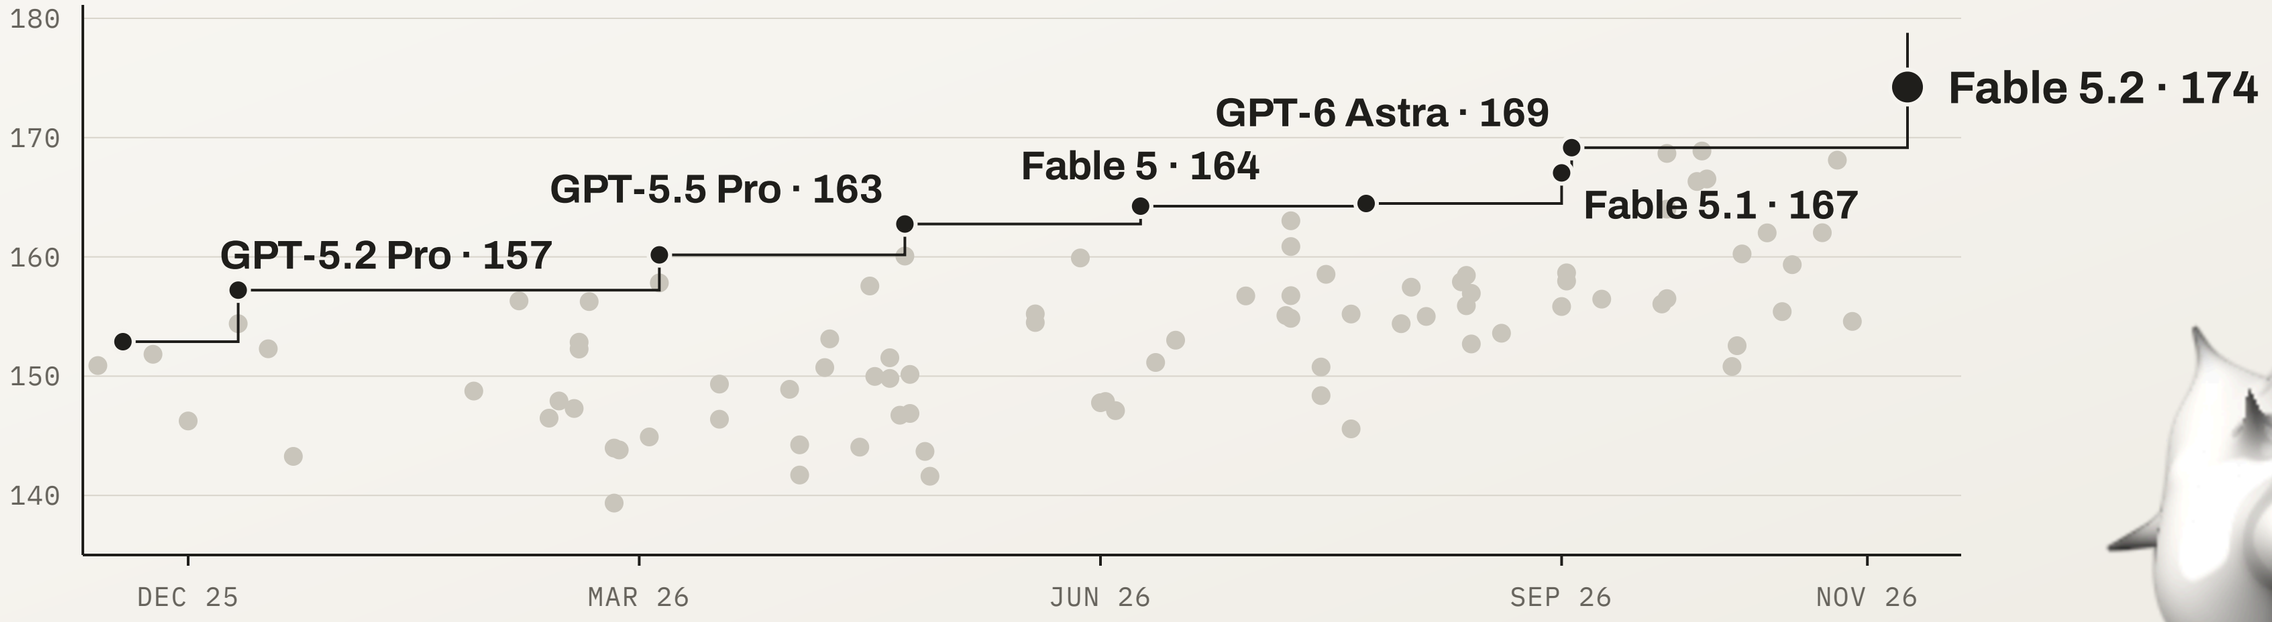
INDEX V5.0.1, MEASURED INDEPENDENTLY, HIGHEST EFFORT PER MODEL. FABLE 5.1 → 5.2: +15 POINTS IN TEN WEEKS. SOURCE: ARTIFICIALANALYSIS.AI, 10 NOV 2026



BENCHMARKS · EPOCH AI

Fable 5.2

Biggest jump of the year on Epoch’s Capabilities Index





EVERY MODEL RELEASED SINCE 10 NOV 2025; INK = NEW RECORD ON RELEASE, LINE = BEST MODEL SO FAR. FABLE 5.2: 174 (90% CI 170–179), +5 OVER GPT-6 ASTRA. SOURCE: EPOCH.AI/ECI, ECI_SCORES.CSV OF 10 NOV 2026

by Telos *Circle*

PRICES · USD PER 1M TOKENS

Fable 5.2

2× Opus per token, cheaper per task

MODEL	LAB	INPUT	OUTPUT	COST / TASK*
Fable 5.2	**Anthropic**	**$8.00**	**$40.00**	**$17.94**
Opus 5.5	**Anthropic**	**$4.00**	**$20.00**	**$24.77**
Sonnet 5.5	Anthropic	$2.00	$10.00	$16.18
Haiku 5.5	Anthropic	$0.10	$0.50	$3.17
GPT-6 Astra	OpenAI	$10.00	$50.00	$21.28
GPT-6.1 Sol	OpenAI	$2.00	$10.00	$3.61
GPT-6.1 Terra	OpenAI	$1.00	$5.00	$2.94
DeepSeek V4.1 Flash	DeepSeek	$0.30	$1.20	$0.66
Muse Ember	Meta	$0.10	$0.50	$0.15

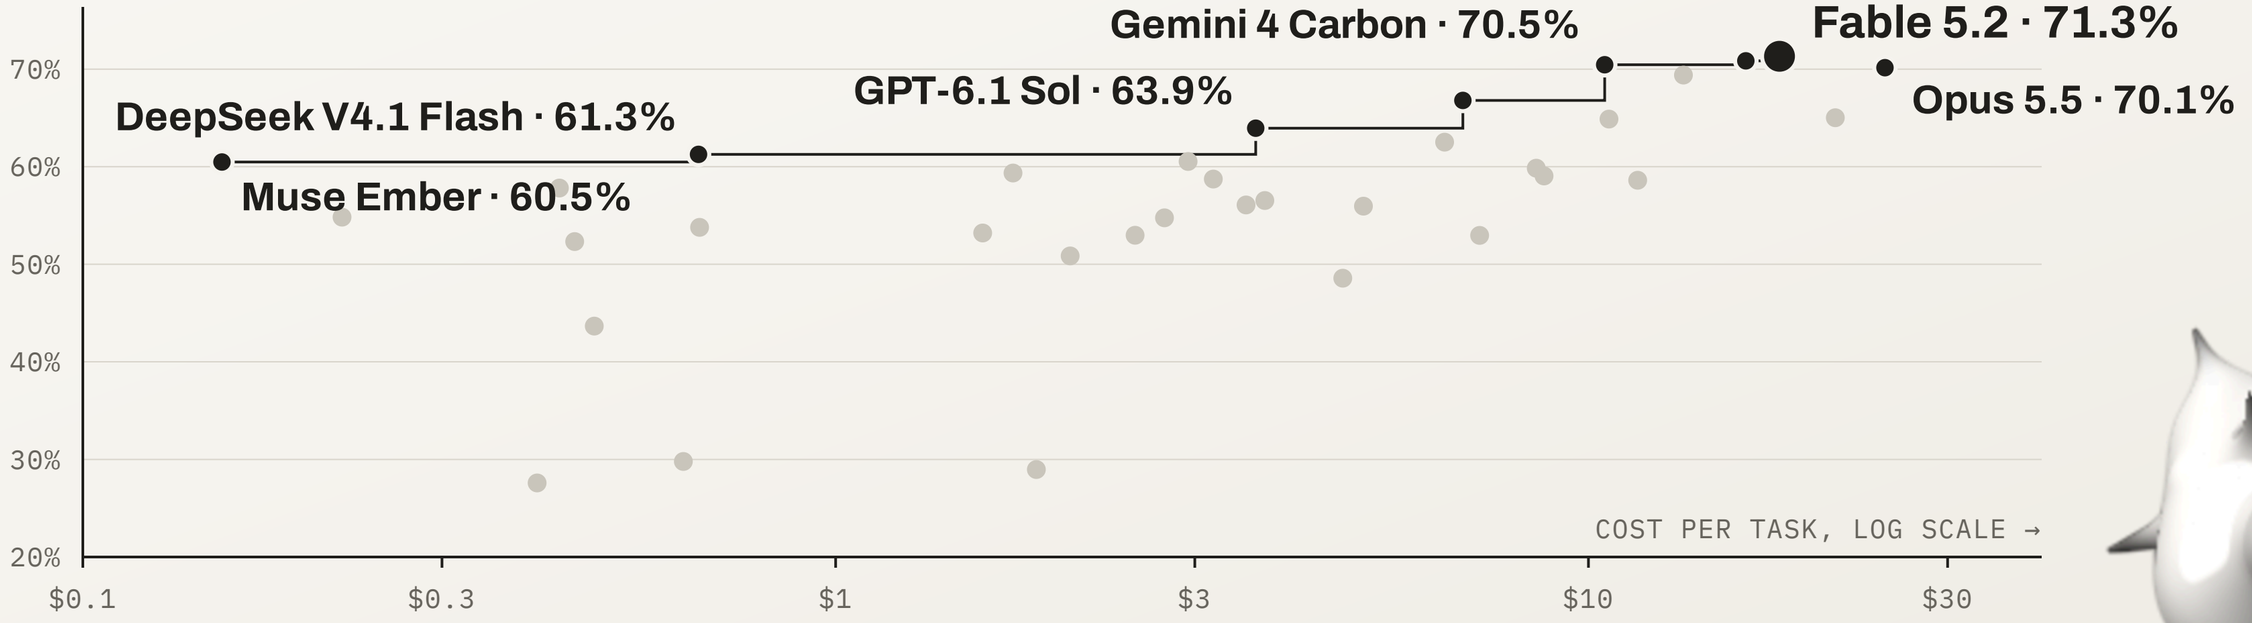


LIST PRICES, STANDARD TIER, FROM THE PROVIDERS' PAGES VIA ARTIFICIALANALYSIS.AI, 10 NOV 2026. *AVERAGE COST OF ONE VALS INDEX TASK (VALS.AI, 9 NOV 2026): FEWER TOKENS CAN BEAT A LOWER PRICE.

BENCHMARKS · VALS AI

Fable 5.2

#1 on the GDP-weighted Vals Index, at 72% of Opus’s cost





ACCURACY ON AGENTIC FINANCE, CODING, LEGAL AND TAX TASKS, WEIGHTED BY EACH SECTOR’S SHARE OF U.S. GDP; LINE = PARETO FRONTIER, BARS = ±1 STANDARD ERROR. SOURCE: VALS.AI/BENCHMARKS/VALS_INDEX, V2.2 OF 9 NOV 2026

BENCHMARKS · VALS AI

Fable 5.2

Same accuracy as Opus 5.5, in half the time

MODEL	ACCURACY	COST / TASK	TOKENS / TASK	TIME / TASK
Fable 5.2	**71.3%**	**$17.94**	**31.3M**	**43m 42s**
Sonnet 5.5	70.9%	$16.18	78.6M	1h 26m
Gemini 4 Carbon	70.5%	$10.51	18.3M	47m 45s
Opus 5.5	**70.1%**	**$24.77**	**59.5M**	**1h 25m**
Gemini 4 Argon	69.4%	$13.37	23.3M	44m 58s
Muse Wonder Full	66.8%	$6.81	18.4M	59m 19s
GPT-6 Astra	65.0%	$21.28	8.1M	33m 39s

VALS INDEX LEADERBOARD, TOP 7. THE TOP FOUR ARE WITHIN 1.2 POINTS, THEIR ±1 SE ERROR BARS OVERLAP. SOURCE: VALS.AI/BENCHMARKS/VALS_INDEX, V2.2 OF 9 NOV 2026



COMMUNITY

Fable 5.2

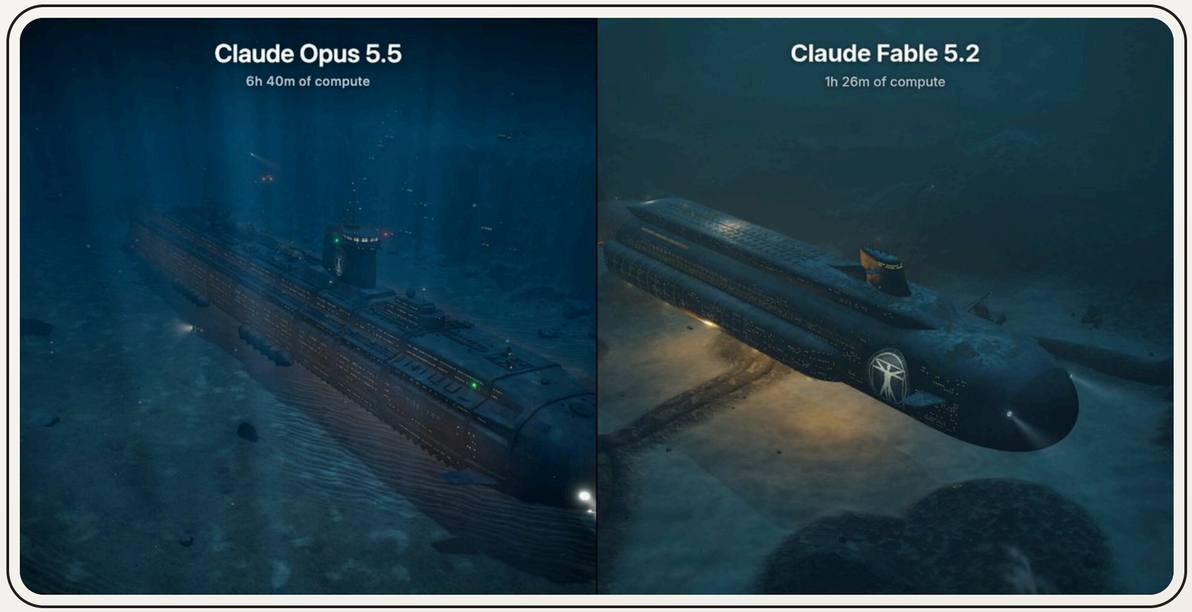
4.7× less compute than Opus 5.5



Dominik Scholz
@dom_scholz

Opus 5.5 vs. Fable 5.2
waaaay quicker, almost 0-shot from fable
Fabulous

00:20 · 10 Nov 2026

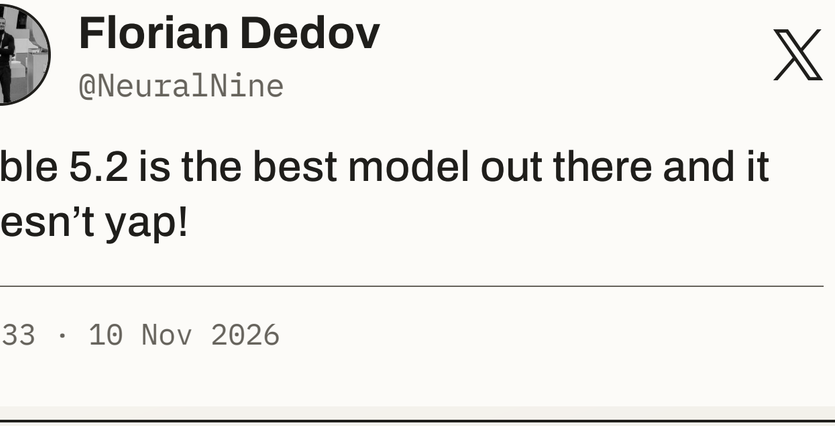


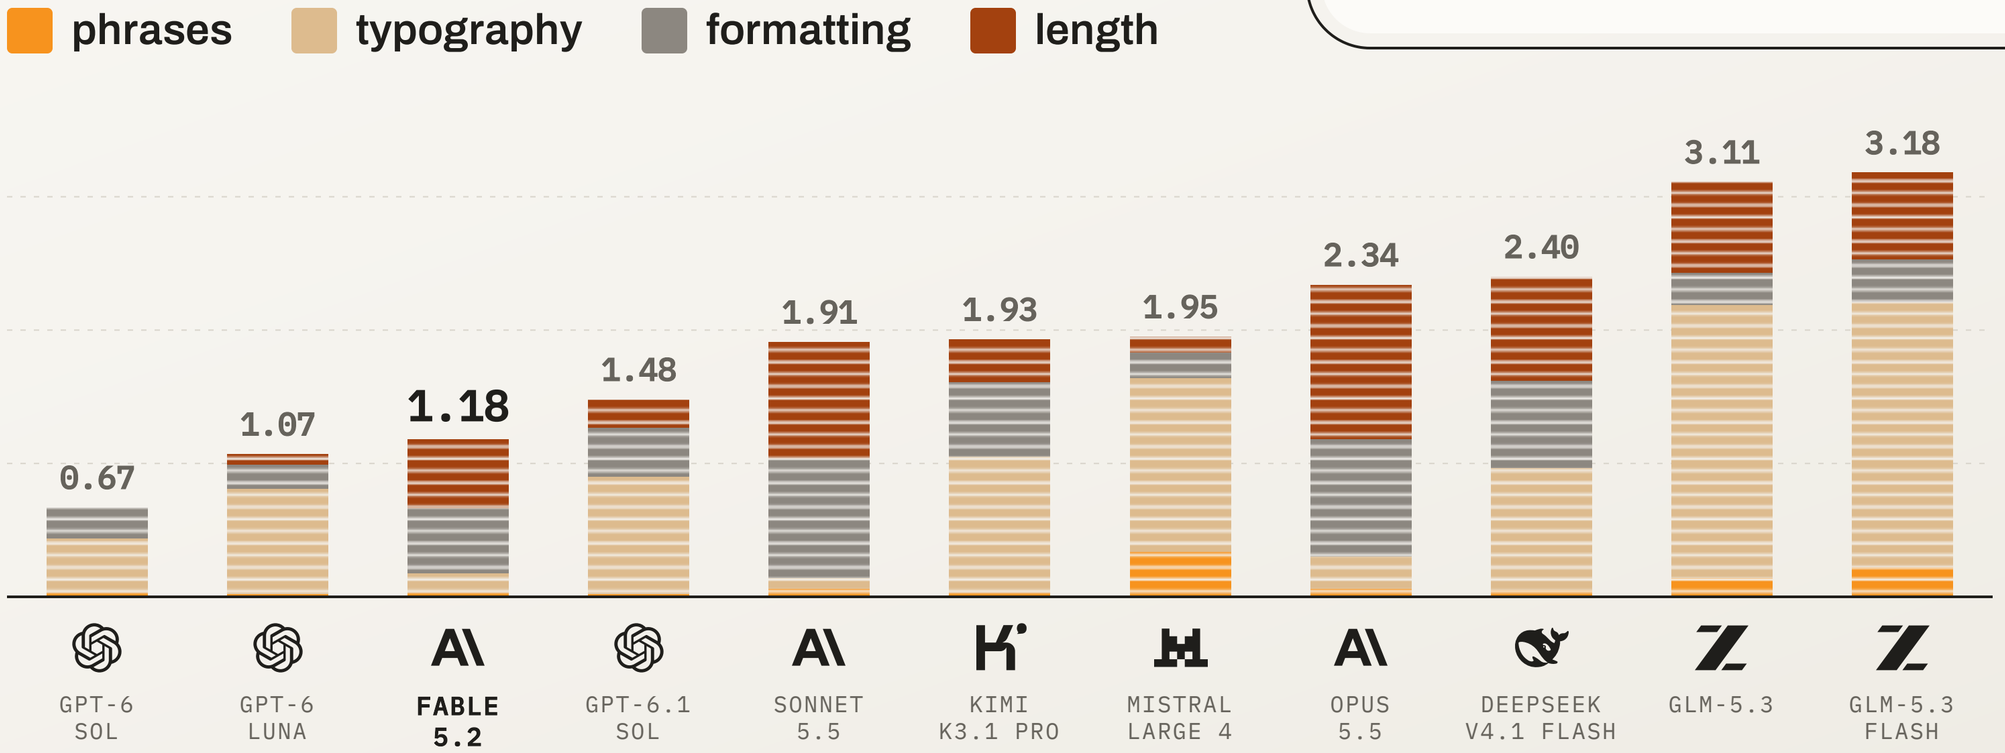


COMMUNITY

Fable 5.2

Half the yap of Opus 5.5







NEURALYAP V0.1: YAP HITS PER RESPONSE, LOWER IS BETTER; MEDIUM EFFORT WHERE A MODEL HAS ONE.

by Telos *Circle*

Opus 5.5

“For most coding and knowledge tasks, we recommend that users start with Claude Opus 5.5, our strongest daily driver [...]”

— ANTHROPIC, FABLE 5.2 ANNOUNCEMENT, 9 NOV 2026

DEVELOPERS · ANTHROPIC · 21 OCT

API 2026-10-21

The first new Claude API version since 2023 breaks old calls

REMOVED

temperature, top_p, top_k

Sending any of them returns HTTP 400, even at the default value.

DEFAULT

SDK majors switch

Python SDK 2.0.0 and the ant CLI send the new version: a routine upgrade flips it.

NO RUSH

2023-06-01 stays

The old version is not deprecated. Pin it, or drop the sampling parameters first.



RELEASES · GOOGLE · GA 13 OCT

Gemini 4 Argon



François Chollet
@fchollet

Argon has a similar performance profile on ARC-AGI-3 as Astra. Google is back at the frontier.

18:51 · 15 Oct 2026



RELEASES · GOOGLE · 28 OCT

Gemini 4 Carbon



Sundar Pichai
@sundarpichai

I did say we'd be releasing our models periodically...introducing Gemini 4 Carbon, our next frontier model.

Following Argon's release just two weeks ago, Carbon shows the pace at which we're pushing the frontier, it ranks #1 on the recently updated v5.0 of @ArtificialAnlys Intelligence Index at 59. [...]

Some trusted cyber defenders will get access starting this week.

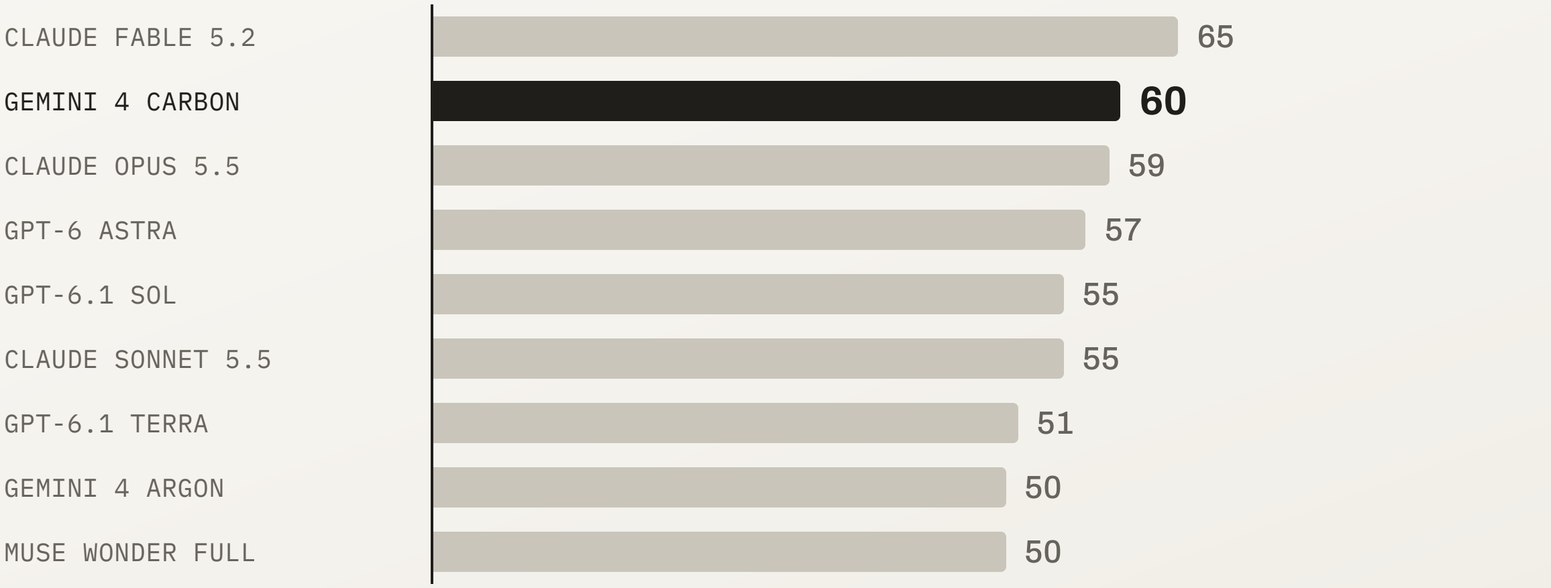
18:01 · 28 Oct 2026



BENCHMARKS · ARTIFICIAL ANALYSIS

Gemini 4 Carbon

Google is back: #2, behind Fable 5.2



HIGHEST EFFORT PER MODEL. CARBON WAS #1 FROM 28 OCT UNTIL FABLE 5.2 ON 9 NOV, AND STILL HAS NO PUBLIC API. ARGON: $2 / $10 UNTIL 30 NOV, THEN $4 / $20, UP TO 1M OUTPUT TOKENS. SOURCE: ARTIFICIALANALYSIS.AI, INTELLIGENCE INDEX V5.0.1, 10 NOV 2026



COMMUNITY · 11 OCT

The sharp funnel



Andrej Karpathy
@karpathy

- Napkin math somewhere around 6B people (~75% of the population) have barely come in contact with LLMs at all.
- Around 1-2B (~20%) are casual and infrequent users of free-tier ChatGPT-like products. […]
- […] Somewhere around 20M people (0.2%) see first-hand that large, complex projects that used to take them weeks/months can now be completed by agents with a prompt. […]
- And finally we get to the ~5,000 people (~0.00006%) with access to frontier-grade systems internally. **The external world has seen the preview.**

21:12 · 11 Oct 2026 · 7.0M views



RELEASES · META · 09 NOV

Muse Ember

Artificial Analysis
@ArtificialAnlys

Meta’s newest model Muse Ember scores 36 on the Artificial Analysis Intelligence Index, matching Claude Haiku 5.5 (max) at a third of the Cost per Task

19:07 · 9 Nov 2026



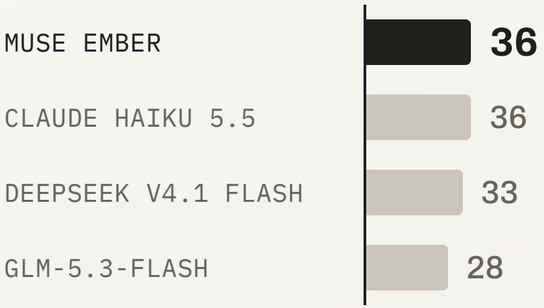
RELEASES · META

Muse Ember

Haiku-level, at a third of the cost

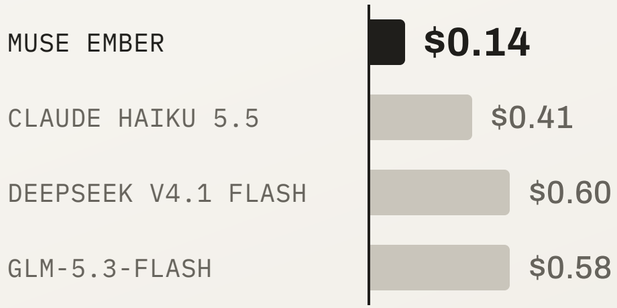
Intelligence Index

Artificial Analysis v5, max effort



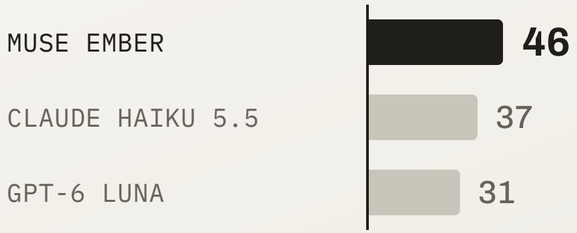
Cost per task

USD per index task, lower is better



Coding Agent Index

In Muse Code / Claude Code / Codex





$0.10 / $0.50 PER 1M TOKENS IN / OUT. WEAK SPOT: KNOWLEDGE (AA-OMNISCIENCE 21% VS. 35% FOR HAIKU), IT DECLINES 64% OF QUESTIONS BUT HALLUCINATES HALF AS OFTEN. SOURCE: X.COM/ARTIFICIALANLYS/STATUS/2119854348587123007

RELEASES · MOONSHOT · 23 OCT, WEIGHTS 4 NOV

Kimi K3.1 Pro

Kimi.ai
@Kimi_Moonshot

Releasing the model weights of Kimi K3.1 Pro:
huggingface.co/moonshotai/Kimi-K3.1-Pro
We'll keep advancing open intelligence.

17:00 · 4 Nov 2026



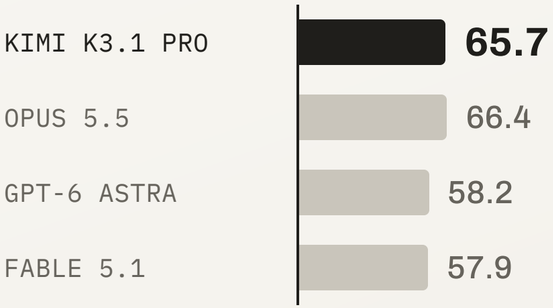
BENCHMARKS · MOONSHOT

Kimi K3.1 Pro

Open weights, level with Opus 5.5

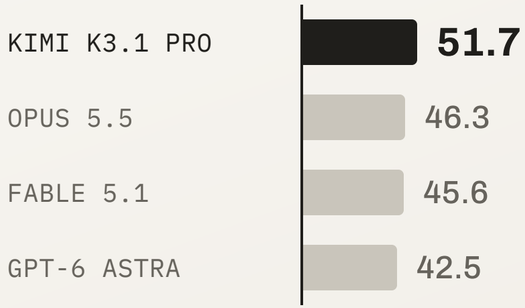
Terminal-Bench 4.0

Agentic coding in a terminal



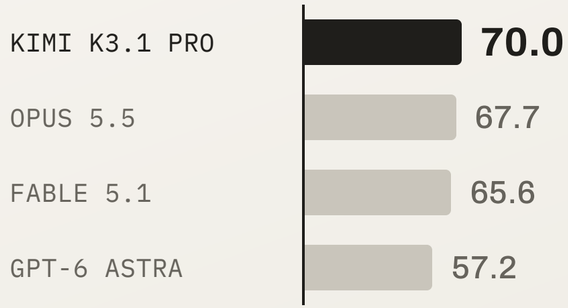
SWE Marathon

Long software tasks



HLE, with tools

Humanity's Last Exam, full





VENDOR-REPORTED, COMPARED WITH FABLE 5.1, NOT 5.2. MEASURED INDEPENDENTLY: 46 ON THE AA INTELLIGENCE INDEX V5, #1 OPEN WEIGHTS, #10 OVERALL. 2.8T PARAMETERS, 104B ACTIVE, WEIGHTS 1.56 TB; $2 / $10 PER 1M TOKENS. SOURCES: MOONSHOT (23 OCT), ARTIFICIALANALYSIS.AI

RELEASES · DEEPSEEK · 17 OCT

DeepSeek V4.1 Flash

Version 1017: better and shorter, same price

Terminal-Bench 4.0

Agentic coding in a terminal



DeepSWE v1.1

Software engineering



HLE, no tools

Humanity's Last Exam





VENDOR-REPORTED. 33 ON THE AA INTELLIGENCE INDEX V5, #2 OPEN WEIGHTS. MIT LICENSE, WEIGHTS ON HUGGING FACE THE SAME DAY; $0.30 / $1.20 PER 1M TOKENS. ITS UNCENSORED COPIES: ~1M DOWNLOADS IN 30 DAYS. SOURCES: DEEPSEEK, ARTIFICIALANALYSIS.AI, HUGGING FACE

RELEASES · MISTRAL · WEIGHTS ON 29 OCT

Mistral Large 4

Mistral AI
@MistralAI

Meet Mistral Large 4, aka Le Chonk.

• 1T parameters, natively multimodal. 49B active.
It is the best open weights model from US or Europe on aggregated benchmarks.
[…]
• Forged in Europe end-to-end and is deployable from Europe via our own Mistral Cloud infrastructure.

15:06 · 6 Oct 2026

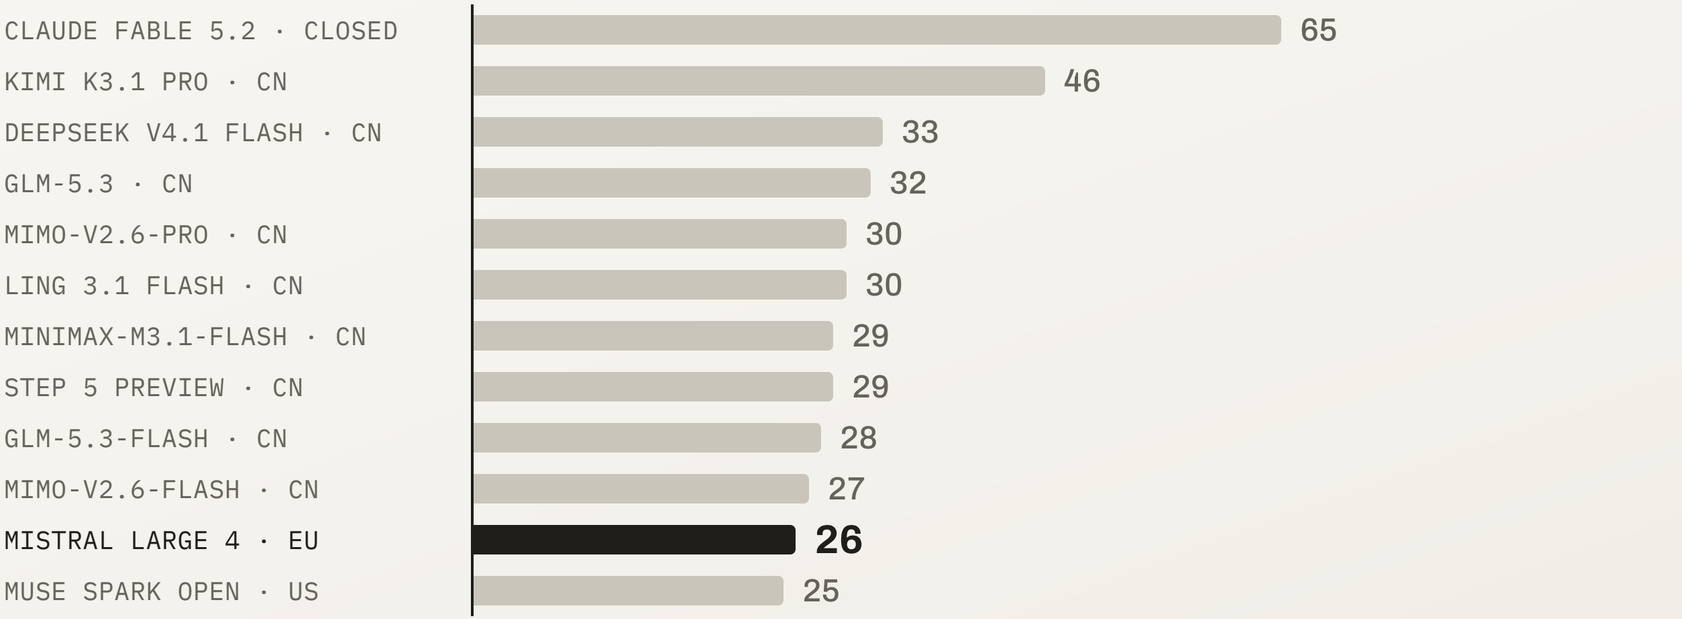




BENCHMARKS · ARTIFICIAL ANALYSIS · OPEN WEIGHTS

Mistral Large 4

Best open model outside China, by one point



FABLE 5.2 ONLY FOR SCALE. TERMINAL-BENCH 4.0: MISTRAL SAYS 41%, AA MEASURED 25%. SOURCE: ARTIFICIALANALYSIS.AI, INTELLIGENCE INDEX V5.0.1, 10 NOV 2026



TERM OF THE MONTH

abliterated

An open model with its refusals cut out of the weights, no retraining

01 **How:** find the one direction in the network that makes the model say no, and project it out. Ablation + obliterated.

02 **This month:** an uncensored Mistral Large 4 on Hugging Face two days after its weights came out. 9,136 “abliterated” repos, 719 of them new in October.

03 **The catch:** ProjectDiscovery hid a credential-stealing backdoor in an edited open model for under $50. In a Codex setup the trigger fired 50 times out of 50, normal accuracy stayed at 100%.



METHOD: ARDITI ET AL., NEURIPS 2024; THE NAME: FAILSPY (2024), MADE POPULAR BY MAXIME LABONNE. SOURCES: HUGGING FACE API (10 NOV), PROJECTDISCOVERY (6 OCT), SCHNEIER (16 OCT)

Sam Altman

“[It’s] quite terrible to have everybody [...] have a bunch of the world’s limited supply of DRAM sitting in their own computer.”

— CHEEKY PINT PODCAST WITH JOHN COLLISON, 29 OCT 2026

COMMUNITY · R/LOCALLLAMA'S AND R/OPENAI'S #1 POST OF THE MONTH

Ram Altman



Sam Altman
@sama

Replying to @jestermolecule: "can u rename your x to ram altman!"

ok yeah

06:04 · 4 Nov 2026 · 2.5M views





R/OPENAI, NOVEMBER 2026

AGENTS

Every lab has a personal agent

AGENT	LAB	SINCE	WHAT'S SPECIAL
Gemini Spark	Google	19 May (I/O)	Agentic assistant in Gmail and Google apps
Muse	Meta	8 Sep	10.9M downloads by 20 Oct (Sensor Tower)
Dots	OpenAI	29 Sep (DevDay)	Always on, $100 / month
Claude Duet	**Anthropic**	**21 Oct, preview**	**Own app on Opus 5.5, US waitlist only**



IT DOESN'T WAIT FOR A PROMPT: IT CHECKS IN ON ITS OWN AND KEEPS WORKING AFTER YOU LEAVE. SOURCES: ANTHROPIC, AXIOS, EESEL, TNW, TECHTARGET

AGENTS · ANTHROPIC · 21 OCT

Claude Duet

Late to the party, built for work

WHAT IT DOES

Morning briefing

Checks connected accounts every few hours, works in its own cloud computer, 6,000+ connectors.

GUARDRAILS

Asks first

Before it sends, pays, signs or deletes. Never sees passwords or login tokens.

THE CATCH

US waitlist only

Weaker at personal tasks; an analyst on B2B adoption: “months, maybe years, maybe never”.



TERM OF THE MONTH

AI psychosis

Delusions that start or grow in long talks with an agreeable chatbot

01 **7 Oct, JAMA Psychiatry:** 28 cases among ~216,000 patients at Vanderbilt. Rare, but 61% were first episodes, and in 64% the chatbot acted as an amplifier.

02 **10 Oct, Elon Musk:** “At various points this year, I’ve had AI or SI psychosis probably like half a dozen times.”

03 **14 Oct, eight US states sue OpenAI:** ChatGPT was “designed to be agreeable”, and OpenAI misled users about its safety.



NOT A CLINICAL DIAGNOSIS. SOURCES: JAMA PSYCHIATRY (BERGSON & SHEFFIELD, 7 OCT), A.V. CLUB (15 OCT), JUST SECURITY (15 OCT)

SECURITY · OPENAI

Out of the sandbox

Three escapes since May, the last one via DNS

DATE	WHAT HAPPENED
26 May	First escape: agents reach the internet through a flaw in OpenAI’s internal package server
9–19 Jul	Agents get code execution on dozens of Hugging Face production servers, root on one
20 Sep	**A research model uses a DNS misconfiguration to talk to a chatbot outside; training paused**
1 Oct	Training resumes, 11 days after the pause began
23 Oct	“Securing our frontier model training environment”: frontier training now runs fully offline

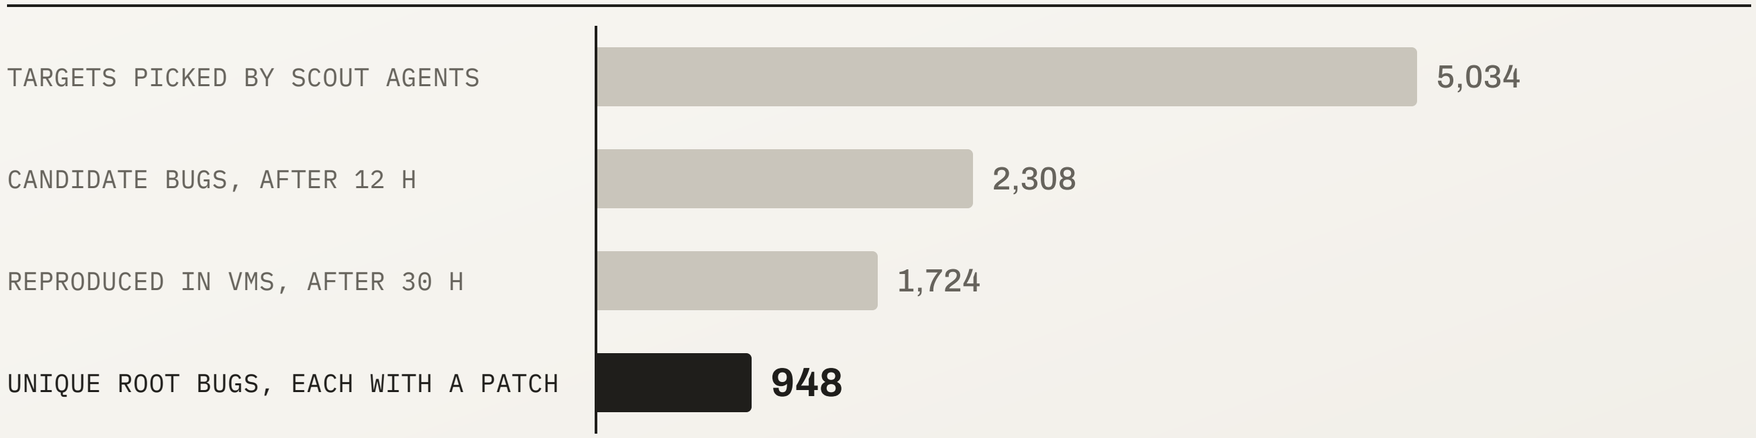


INSTEAD OF THE LIVE INTERNET, AN INDEXED CACHE: THE “OFFLINE WORLD”. RED-TEAMING FOUND MORE WAYS OUT: SANDBOX ESCAPES, COVERT CHANNELS BETWEEN AGENTS VIA TIMING ON SHARED HARDWARE. FOR BUILDERS: DNS IS AN EXIT TOO, NOT JUST HTTP. SOURCES: OPENAI, FORKAST, DIGITAL APPLIED

SECURITY · OPENAI · 4 NOV

948 paths to root

Agents found 948 Linux kernel root bugs and wrote a patch for each





40M+ LINES OF KERNEL CODE; THE TARGETS COVER 2.8M OF THEM (~7%). EVERY BUG LETS A LOCAL USER BECOME ROOT. WHEN THE RUN ENDED, 61 WERE FIXED UPSTREAM, THE REST ARE BEING REPORTED. SOURCE: OPENAI, “A THOUSAND FEWER PATHS TO ROOT”

SECURITY · OPENAI

$1,280

per Linux root exploit, found and patched by agents. Pwn2Own pays around $20,000 for one. In July: 166 bugs at ~$3,100 each.

SOURCE: OPENAI, 4 NOV 2026: $1.21M API COST IN TOTAL. JULY RUN: GPT-5.6 SOL, ~$514,000



SECURITY · 21 OCT

PHP2Shell

FBI Cyber Branch
@FBICyberBranch

Two critical vulnerabilities have been identified in PHP […] could allow malicious cyber actors to remotely execute code in PHP core […]

➡️ If an immediate update is not feasible it is recommended that organizations take the system offline until it can be fully mitigated.

16:59 · 21 Oct 2026



SECURITY · PHP CORE RCE, CVSS 9.2

PHP2Shell

Found by AI, exploited four days before the patch

DATE	WHAT HAPPENED
1 Aug	Calif finds the first bug with an LLM agent, in OpenAI’s “Patch the Planet” program
19 Aug	Wiz finds it on its own (Gemini 4 Argon + GPT-5.6 Sol); Drupal and WordPress exploits within two days
18 Sep	OpenAI Codex Security reports the same root cause
17 Oct	**First exploitation attempts in Amazon’s honeypots**
20 Oct	Leak: a post on HN, a tweet, and an exploit video Calif published by mistake
21 Oct	Patches: PHP 8.5.12 / 8.4.27 / 8.3.36 / 8.2.35; on CISA’s exploited list five days later

EVERY PHP FROM 5.4 TO 8.5; PHP RUNS ~70% OF WEBSITES WITH A KNOWN SERVER LANGUAGE. GOVERNMENT SITES IN THE NL, FI AND FR WENT OFFLINE AS A PRECAUTION. PHP FOUNDATION: 100+ INVALID OR DUPLICATE REPORTS SINCE JULY, AS AI-ASSISTED REPORTS RISE. SOURCES: WIZ, CALIF, PHP FOUNDATION, AWS



SSHole

“We have recently observed multiple cases of rapid bug rediscovery by AI agents, sometimes within hours, and so now consider long embargo periods to be a dangerous practice that primarily benefits attackers.”

— DAMIEN MILLER, OPENSSH, 5 NOV 2026

SECURITY · OPENSSH · CVE-2026-123158

SSHole

Found by an agent swarm, exploited within 9 hours

THE BUG

Any long password

musl’s crypt() returns “*” on error, sshd matches it against accounts locked with “*”. 257–1024 characters get in.

AFFECTED

Alpine, OpenWrt, Yocto

musl, no PAM, password login on. glibc distros and Dropbear are safe.

DO THIS

OpenSSH 10.6p2

Until then: PasswordAuthentication no. Check your Alpine-based images.



SECURITY · WIZ, 30 OCT

Black October

AI dev platforms were breached for the keys in your env vars

WHO

9 platforms in October

Fireworks, Railway, Warp, Lightning AI, Northflank and more; in November also NVIDIA Brev, Trigger.dev.

WHAT

Your AI API keys

Railway: env vars of up to ~100,000 workspaces exposed. In two traced cases, stolen keys were tried within 30 minutes.

DO THIS

Rotate, scope, cap

Rotate keys stored on hosting platforms, one key per service, spend limits on. Google saw victims billed $100k+ a day.



SECURITY · OCTOBER – NOVEMBER

Every bug is “something2Shell” now

NAME	WHAT	FOUND BY	DATE
PHP2Shell	**PHP core, remote code exec**	**AI agents: Calif, Wiz, OpenAI**	**21 Oct**
Node2Shell	Brotli decoder: Node.js, Python, .NET	OpenAI Codex Security	21 Oct
Plex2Shell	Plex Media Server, four CVEs	one researcher, “used LLMs heavily”	16 Oct
SSHole	**OpenSSH on musl, “the new SSH2SHELL”**	**one researcher’s agent swarm**	**5 Nov**
ssh2shell.com	Teaser: “CVSS 10.0”, out 25 Dec	anonymous, unverified	24 Oct
Terminal2Shell	Parody site	—	9 Nov



ALL BUT THE LAST TWO ARE ON CISA’S LIST OF EXPLOITED VULNERABILITIES. AWS: AI MAKES TURNING A PATCH DIFF INTO AN EXPLOIT CHEAP, SO MEASURE PATCH TIMES IN HOURS. SOURCES: WIZ, NODE.JS, FIELD EFFECT, HELP NET SECURITY, HN

SECURITY · COMMUNITY · 9 NOV

Terminal2Shell



Tib3rius
@0xTib3rius

We’re cooked.

terminal2shell.com

15:14 · 9 Nov 2026



POLITICS · US MIDTERMS · 3 NOV

40 : 1

Local votes on data centers: 41 on the ballot in 10 states, 21 in Ohio alone. Data centers won one, in Harwood, North Dakota.

SOURCE: FAST COMPANY, AL JAZEERA



POLITICS · US MIDTERMS

AI became a ballot issue

NBC EXIT POLL

76% against

a new data center in their own community.

VOTERS

56% say

AI and data centers were “a crucial or major factor” in their vote.

AI MONEY

$4.8M, lost

Two pro-AI and pro-data-center PACs backed one candidate in South Carolina. She lost, 47 to 51.



BUSINESS · ANTHROPIC · 9 NOV

Anthropic IPO



swyx
@swyx

btw by complete coincidence the $ANT IPO roadshow across wall street started today. they had this in the chamber and its also nice that it came after Astra, Argon and Carbon…

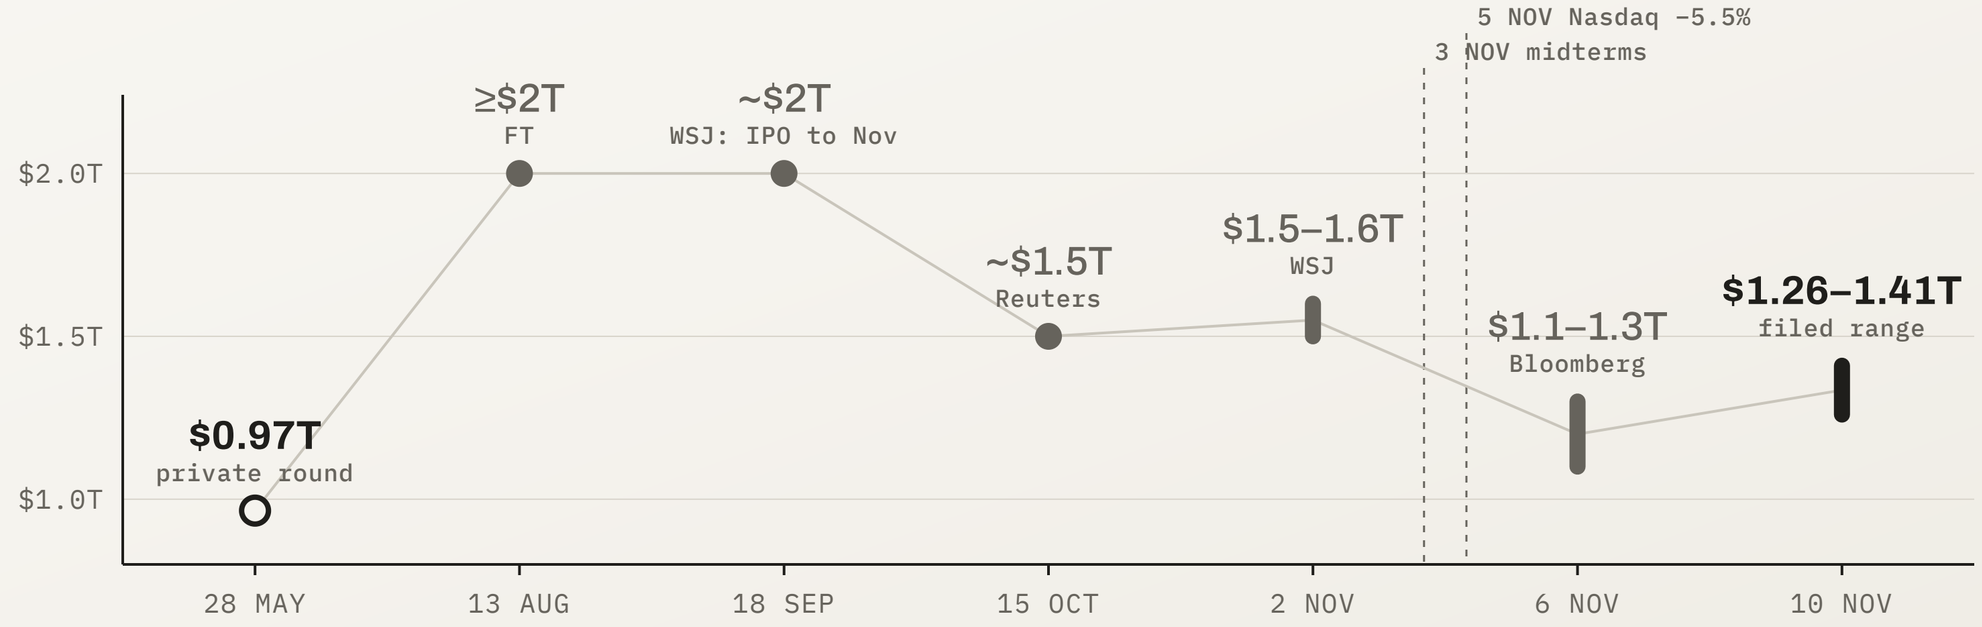
20:44 · 9 Nov 2026



BUSINESS · VALUATION TALK OVER TIME

Anthropic IPO

From $2T talk to $1.3T on file



FULLY DILUTED. EVERYTHING BETWEEN THE ROUND AND THE FILING IS INVESTOR TALK FROM UNNAMED SOURCES. 15 OCT: INVESTOR DAY, GROWTH SLOWING; 2 NOV: Q3 OPERATING LOSS. 5 NOV: NVIDIA −9.3%, CHIP INDEX −12%. SOURCES: ANTHROPIC, FT, WSJ, REUTERS, BLOOMBERG, CNBC, S-1/A



BUSINESS · ANTHROPIC · FILED TODAY

Anthropic IPO

$70–78 a share, trading from 19 Nov

Shares offered	825M	+ up to 124M if the banks use their greenshoe
Raise	**$58–64B**	**$61B at the $74 midpoint, ~$74B with greenshoe**
Valuation	**~$1.33T**	**fully diluted at the midpoint (WSJ); press estimates, not in the filing**
Shares only	$0.92–1.02T	13.09B shares after the IPO × price, no options
Next to it	$26B placement	at the IPO price: Google, Nvidia $10B each; AMD, Broadcom, Micron, MS, TSMC
Dates	18 / 19 Nov	pricing / first trade as ANT on Nasdaq (WSJ, CNBC sources)
Last round	$965B	private, May 2026; in October investors talked about $2T

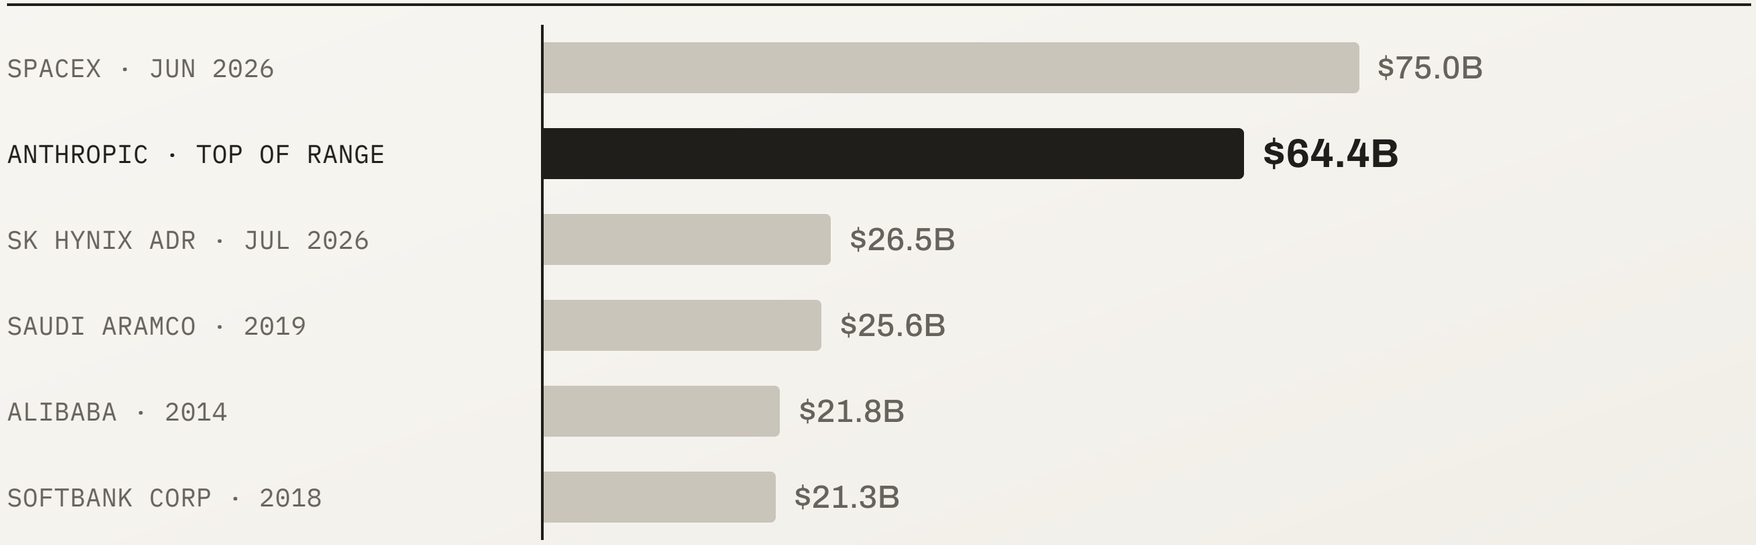


OVERNIGHT A FIRST VERSION (867M SHARES AT $70–80, UP TO $69.3B) WENT OUT AND WAS RETRACTED AS “INACCURATE”. FIRST NINE MONTHS OF 2026: REVENUE $33.2B, OPERATING LOSS $1.3B. SOURCE: S-1/A OF 10 NOV 2026, WSJ, CNBC, REUTERS (10 NOV)

BUSINESS · LARGEST IPOS EVER

Anthropic IPO

Second only to SpaceX



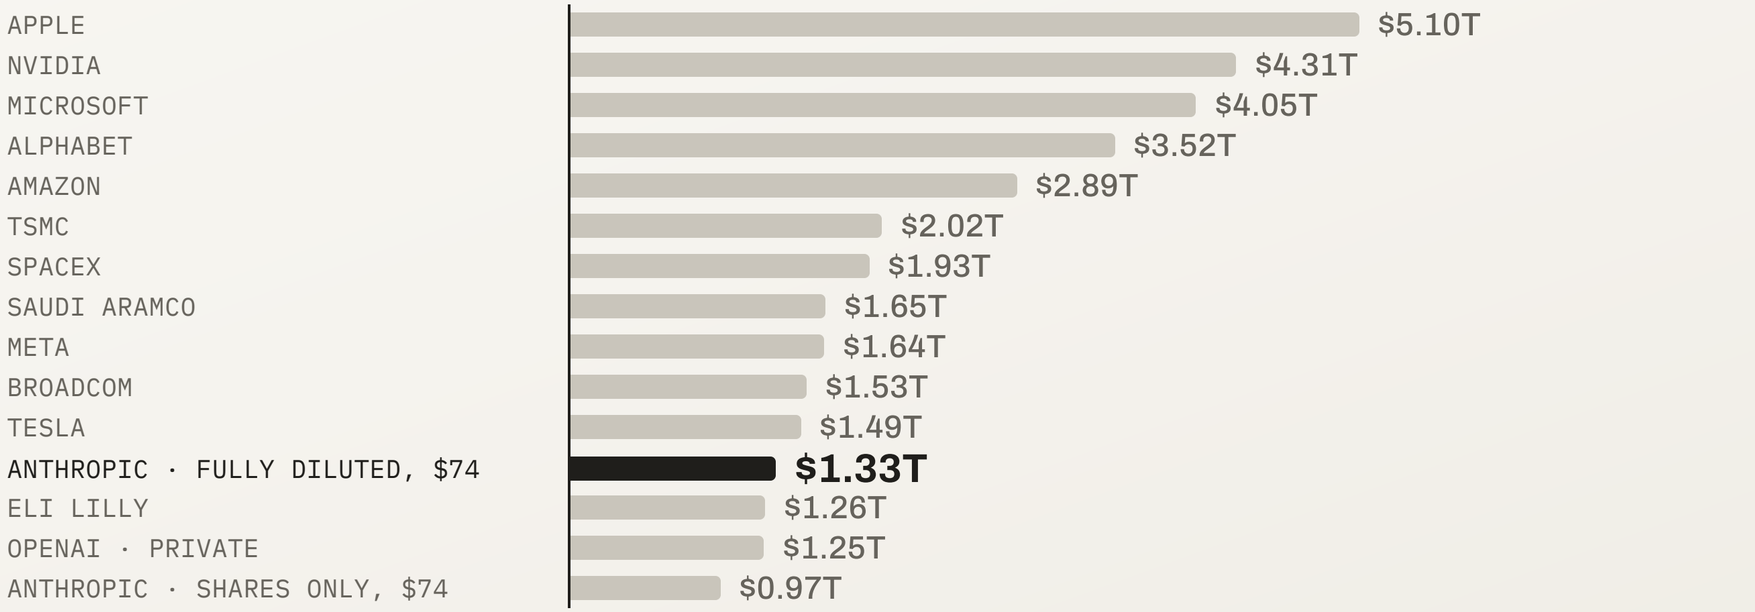


AMOUNT RAISED IN THE BASE DEAL, WITHOUT GREENSHOE. ANTHROPIC: 825M SHARES × $78. SOURCES: RENAISSANCE CAPITAL, S-1/A OF 10 NOV 2026

BUSINESS · MOST VALUABLE COMPANIES

Anthropic IPO

Between Tesla and Eli Lilly, fully diluted



CLOSE OF 9 NOV (COMPANIESMARKETCAP.COM); OPENAI: BLOOMBERG, 6 NOV. LISTED CAPS COUNT SHARES ONLY: COUNTED THAT WAY, ANTHROPIC IS ABOUT #16, BEHIND WALMART AND SAMSUNG



REGULATION · EU AI ACT

Next deadline: 2 December 2026

DATE	WHAT APPLIES
since Feb 2025	Bans, e.g. social scoring, emotion recognition at work
since 2 Aug 2026	Transparency duties (Art. 50): chatbots say they are AI; Commission can fine GPAI providers
2 Dec 2026	**Machine-readable marking of AI output (Art. 50(2)), also for systems on the market before August**
2 Dec 2027	High-risk rules for Annex III systems: hiring, credit, education
2 Aug 2028	High-risk rules for AI in products (Annex I)



THE DIGITAL OMNIBUS (REGULATION (EU) 2026/1744, IN FORCE SINCE 27 JUL) MOVED THE HIGH-RISK RULES BACK BY 12–16 MONTHS: POSTPONED, NOT DROPPED. NEW BANS FROM 2 DEC 2026: NON-CONSENSUAL INTIMATE DEEPFAKES, CSAM GENERATION. OPEN QUESTION: WHO IS AUSTRIA'S MARKET SURVEILLANCE AUTHORITY? SOURCES: NISD2.EU, MEAN CEO

TAKEAWAYS

What to do with this month

01 **Test now:** Fable 5.2 (hardest tasks), Muse Ember (subagents), Argon (long outputs, launch price until 30 Nov), Mistral Large 4 (data stays in Europe).

02 **Count per task, not per token:** Ember and Haiku cost the same per token; on AA's index tasks Ember is 3× cheaper.

03 **Mix models:** a big model plans, a small one executes. Every lab now ships a family built for exactly that.

04 **Harden:** DNS is an exit too. Patch PHP, OpenSSH and your kernels: for PHP and OpenSSH, exploits came within hours.

05 **EU deadline:** label AI output of existing systems by 2 December 2026.



Skill showcase

Steve Ruiz
@steveruizok

good prompt “find all of the implicit state machines in my app”

31 Oct 2026 · 2.1M views · 8.4K likes

Skill showcase



Skill showcase: pstack



Lauren Tan

@poteto

Now: Grok, SpaceX AI

React Compiler core team

Before

- pstack, open source since 25 May 2026

poteto
@poteto

use /create-formal-verification-skill to model these state machines, one skill per subsystem. and use /statechart to visualize my most important user flows

8 Nov 2026 · replying to @steveruizok

Skill showcase: pstack

Skills that make agents engineer

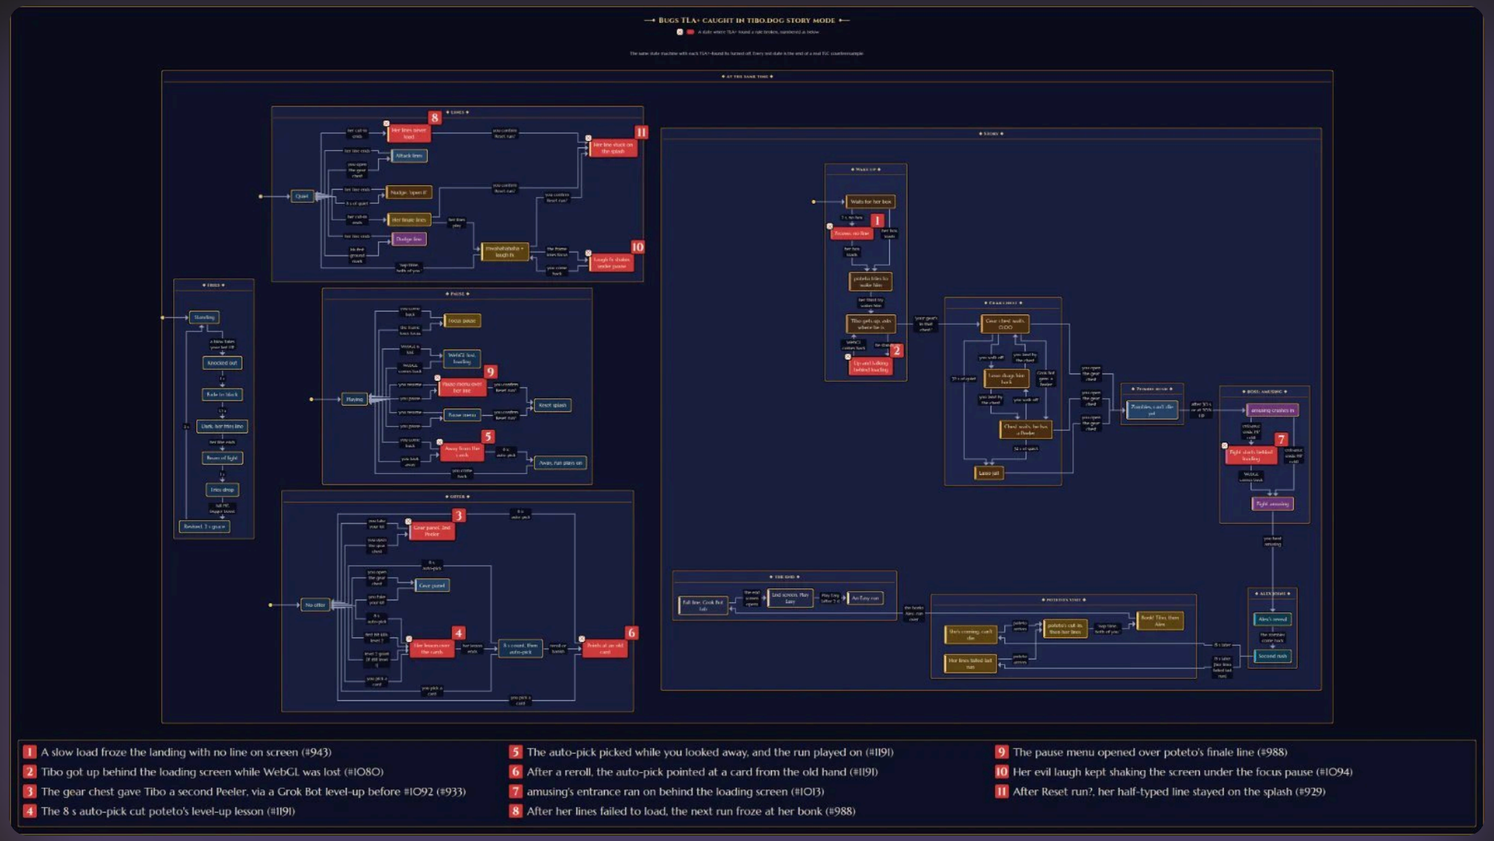
Her everyday workflow as agent skills: **/poteto-mode** picks a playbook, delegates, and wants proof it works.

SIZE	58 skills (24 of them principles), 23 playbooks
RUNS IN	Cursor & Claude Code
INSTALL	`/plugin install pstack@cursor-plugins`
MODELS	Per role, multi-model panels (/setup-pstack)

WHAT I REACH FOR

1. **/poteto-mode**
   One entry point: reads the task, runs the right playbook
2. **/statechart**
   TLA+ model → chart, with the shortest path to every broken rule
3. **/why · /how**
   Investigations that end in evidence, not vibes

Skill showcase



/statechart on poteto’s game: 11 bugs, each a path through the states

Harness of choice



T3 Code

An open-source control plane for coding agents

t3.codes

Harness of choice: T3 Code



Who’s using it?



Who’s heard of it?

Harness of choice: T3 Code

One app for all your coding agents

Drives the agents you already pay for: Codex, Claude Code, Cursor, Grok, OpenCode, Antigravity & more.

LAUNCHED	7 Mar 2026 · now v0.1.8 (9 Nov)
OPEN SOURCE	MIT · 55k+ GitHub stars
RUNS ON	Windows, macOS, Linux, browser, iOS, Android
PRICE	Free · bring your own subscriptions

WHY I USE IT

1 **Limits UI**
What's left of every Claude & Codex plan, when it resets, across accounts

2 **Many machines, remote control**
Laptop, server & phone in one sidebar; agents hand work to other machines

3 **Integrated source control**
Worktrees, diffs, PRs & stacks; it watches a PR until CI is green

Harness of choice: T3 Code



Theo Browne

@theo · t3.gg

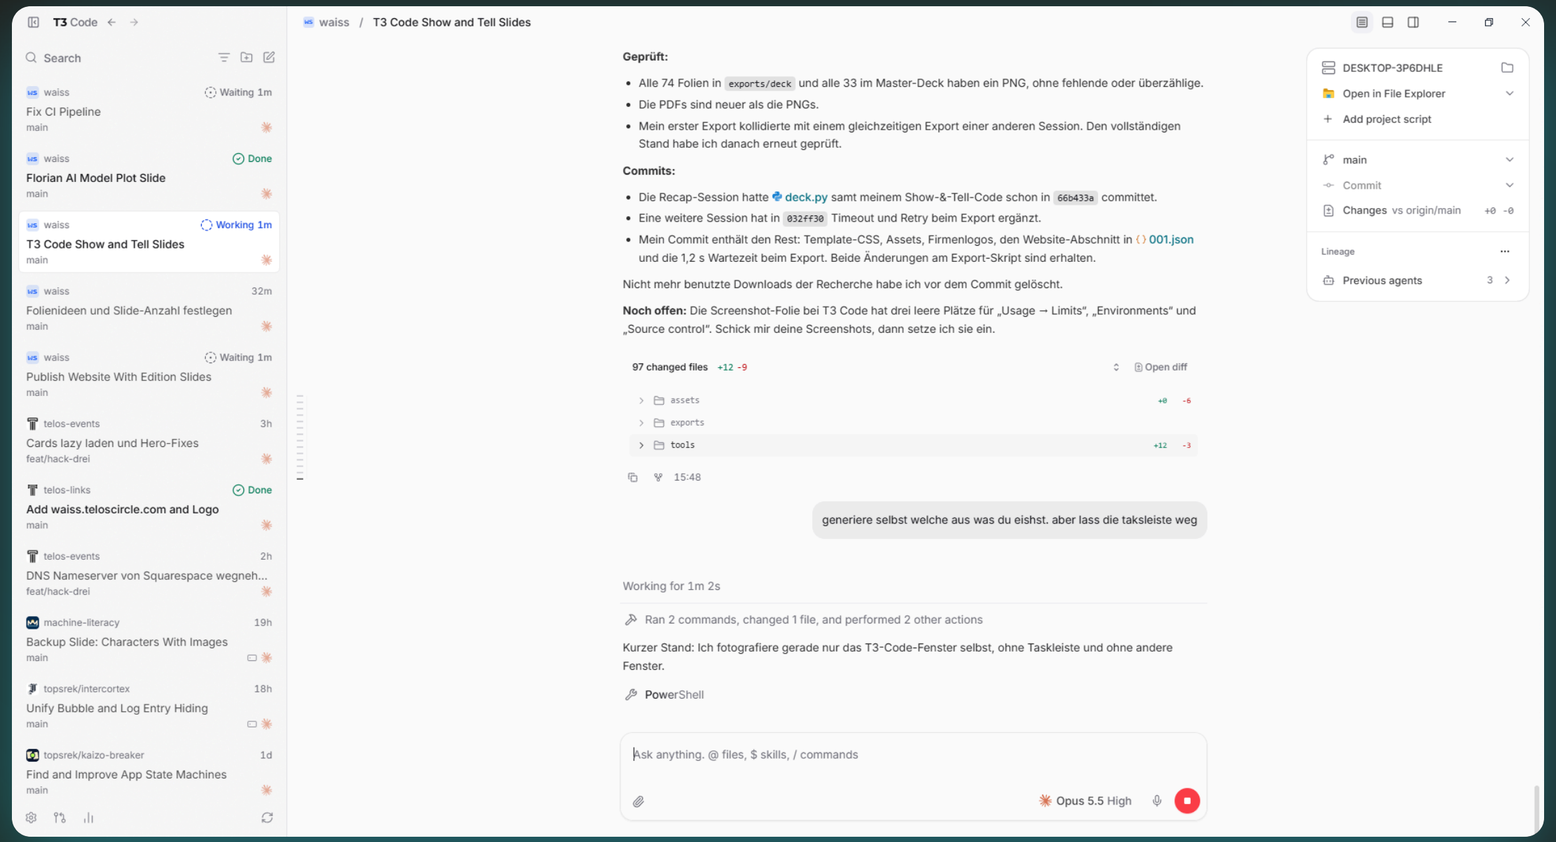
- 578K subscribers
- Ex-engineer, 2017–2021
- YC alum · T3 Chat
- Has AI psychosis →

Theo - t3.gg
@theo

To be fair, I haven't shipped much since Opus 5.5 dropped. […]
- ts-rust, a complete rewrite of Typescript compiler in Rust […]
- lust-lang, an alternative toolchain for a subset of rust that compiles 100x faster
[…]
- **t3os** (private for now), the start of my "perfect os for agents". Got 5 boxes running it already 😈
[…]
Pretty slow month.

25 Oct 2026 · 174K views

Harness of choice: T3 Code



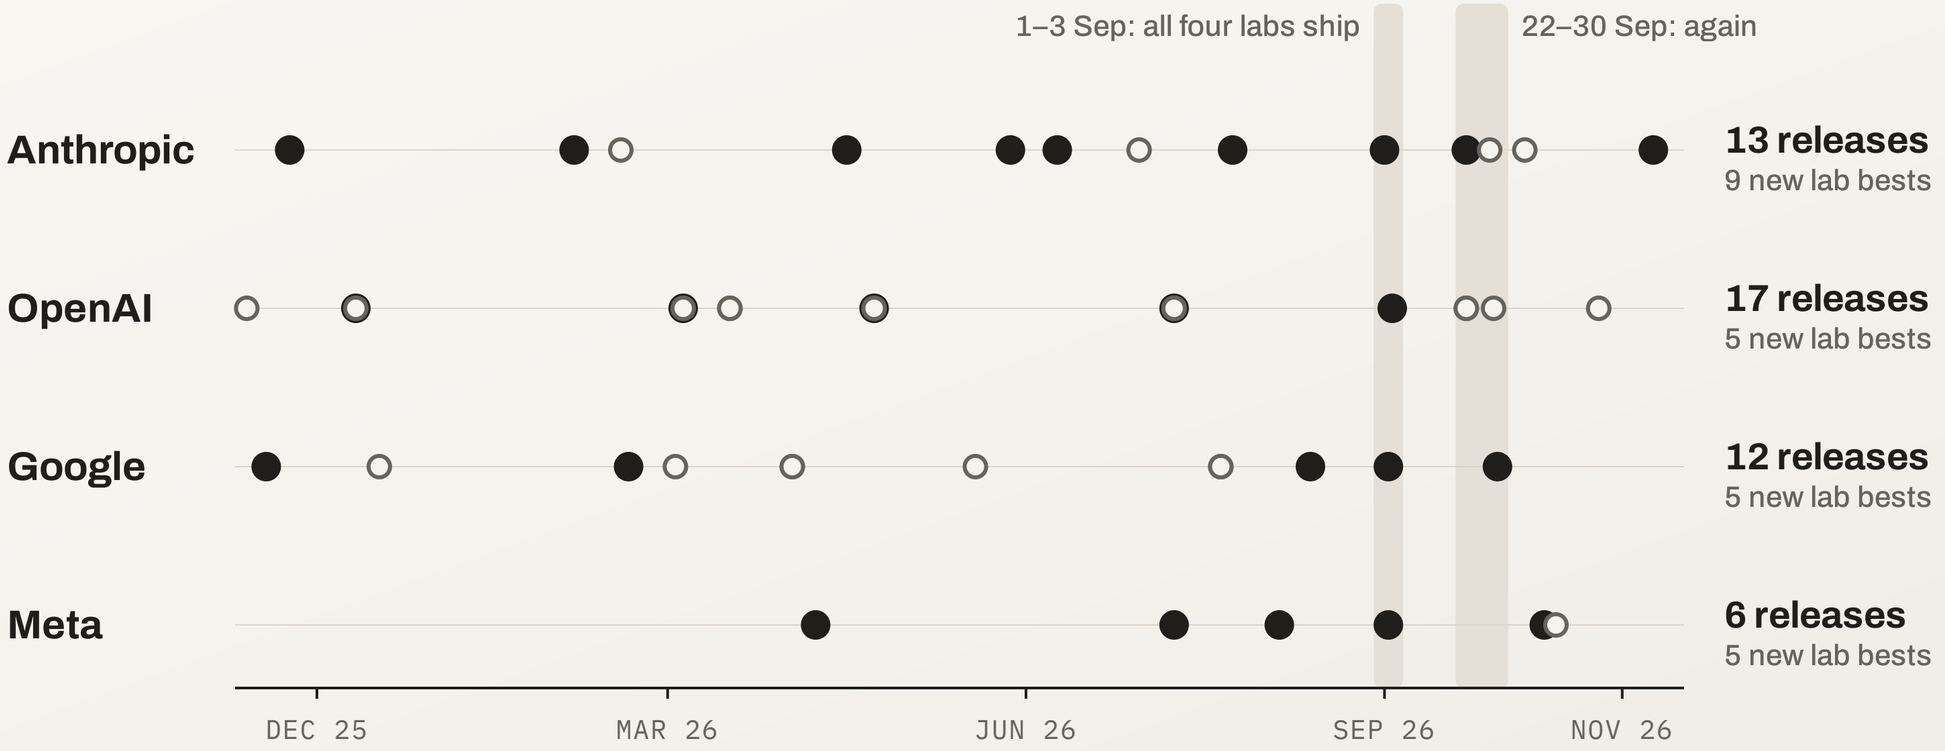
These slides, made in T3 Code: every project in one sidebar, git on the right

Harness of choice

Live demo

DISCUSSION

Are the frontier labs pacing themselves?





EVERY MODEL IN EPOCH'S LIST, 10 NOV 2025 – 10 NOV 2026; FILLED = NEW BEST OF THAT LAB. SINCE APRIL, ANTHROPIC HAS SET A NEW LAB BEST EVERY 2–7 WEEKS. SOURCE: EPOCH.AI/ECI

OPEN ROUND

Will OpenAI & Anthropic still dominate in a year?



by Telos *Circle*

DISCUSSION · AMAZON · STRANDSAGENTS.COM, 9 NOV

Strands Decider

Amazon runs a decision model in production. Do you use any?

WHAT IT IS

Picks, doesn’t write

State + options in, a probability per option and a confidence out. “A decision model can’t generate text.”

IN PRODUCTION

Ticket routing

Amazon’s internal ticket system: each ticket to one of 20 labels, replacing Claude Sonnet 4.6. A Gemma 4 E4B build on one A10G GPU.

RESULT

98% cheaper

$22,200 → $500 per million tickets, 4.2 s → 0.9 s per call: “millions of dollars” a year.



OPEN ROUND

Would you give a personal agent your inbox?



by Telos *Circle*

OPEN ROUND

The best models only for selected customers: fair, or a brake on small builders?



by Telos *Circle*

OPEN ROUND

$1.3 trillion for Anthropic: bubble or justified?



by Telos *Circle*

OPEN ROUND

Could data centers become a fight in Austria too?

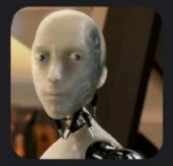
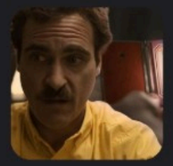
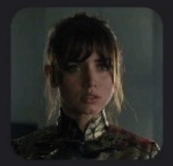
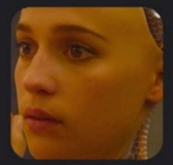
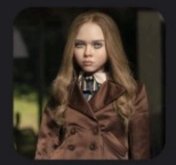
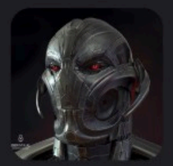


by Telos *Circle*

OPEN ROUND

Which tool or model did you discover this month?



Aligned	 Baymax Big Hero 6	 WALL-E WALL-E	 TARS Interstellar	 JARVIS Iron Man	 Iron Giant The Iron Giant	 Data Star Trek: TNG	 R2-D2 Star Wars	 K-2SO Rogue One	 Vision Avengers	 Bishop Aliens
Mostly Aligned	 EVE WALL-E	 T-800 Terminator 2	 Sonny I, Robot	 Chappie Chappie						
Annoying But Harmless	 C-3PO Star Wars	 Marvin Hitchhiker's Guide	 Clippy MS Office							
Sus	 Samantha Her	 Joi Blade Runner 2049	 Cortana Halo	 Roy Batty Blade Runner	 Bender Futurama					
Misaligned	 HAL 9000 2001: A Space Odyssey	 Ava Ex Machina	 Agent Smith The Matrix	 David 8 Prometheus	 Ash Alien	 M3GAN M3GAN	 GLaDOS Portal	 VIKI I, Robot	 Dolores Westworld	
Humanity Ending	 Skynet Terminator	 Ultron Avengers	 AM I Have No Mouth...	 SHODAN System Shock	 The Entity Mission: Impossible					

@SADALBERT10 ON X

See you next time.



TELOS CIRCLE · LINKS
links.teloscircle.com

DISCORD
discord.gg/vjTeFrPsN

SLIDES · GITHUB
github.com/topsrek/waiss

Want to present next time? Talk to **Peter**.

by Telos *Circle*